

CS-131:HW5 Fun with Sorting

Matthew McLaughlin

May 22, 2018

Contents

1	Shear Sort	2
1.1	Implementation	2
1.2	Analysis	5
1.2.1	Amount of compare exchange operations performed by each proccessor	5
1.2.2	Amount of memory needed by each proccessor	5
1.2.3	Number of Cross-Bar communications cycles	5
2	Bitonic Sort	6
2.1	Implementation	6
2.2	Analysis	9
2.2.1	Amount of compare exchange operations performed by each proccessor	9
2.2.2	Amount of memory needed by each proccessor	9
2.2.3	Number of Cross-Bar communications cycles	9

Assumptions And Notation

I use one based indexing. S is the input sequence $S_{i,j}$ is the subsequence $S_i, \dots S_j$. P_i denotes the i th proccessor in the network. $P_{i,j}$ is the j th value stored in proccessor i . If T is a list of values the $T \Leftarrow x$ adds x to the end of list T . Let M be the name of the network.

1 Shear Sort

For this problem I am assuming that sequence S is of length $\sqrt{n} = k$ where k is positive integer representing the number of processors. This will make it possible to treat the sequence like a square matrix by dividing the sequence into rows and assigning each processor to a row. The assumption that k is a factor of n serves to simplify the pseudocode and analysis. However, it does not affect the complexity of the problem. That is to say, with the assumption that n is a square number and n is divisible by k , we could still in effect treat the sequence as a square matrix as before when $k * k = n$, yet now we would need to simulate this with the processors we had available. For example, suppose $n = 16$ and $k = 2$. In this case we could have P_1 divide its portion of the sequence in half so that the first 4 values and last 4 are treated as two different subsequences allowing P_1 to simulate the effect of having an extra processor.

1.1 Implementation

Below I outline the procedure *Parallel-shear-sort* for performing shear sort on sequence S . This procedure uses helper functions described below. We are required to use a comparison exchange sorting algorithm. Assume that the subroutine *sort* used by *Parallel-shear-sort* is mergesort and that it takes an additional argument specifying whether to sort in increasing order or decreasing order

Algorithm is-even: outputs true if the argument n is even

:

Input: n an integer

Result: true if n is even false otherwise

return $n \bmod 2 = 0$

The crucial part of the *Transpose* routine is line 4 where the i th processor sends its i th row to all of the other processors. The effect is that the old entry i th entry in P_j is overwritten with $P_{i,i}$

Algorithm Distribute-sequence: Stores a subsequence of S in the processors

:

Input: A network of a collection of processor $P_1 \dots P_k$, Sequence S , n is number of values in S

Result: Each processor has a portion of S in its memory. P_i stores the subsequence $S_{a,b}$ where $a = (i - 1)(\sqrt{n}) + 1$ and $b = (i)\sqrt{n}$. Thus, the routine effectively transforms S into a square matrix and distribute it amongst the processors so that P_i stores the i th row of the matrix

```
1  $i \leftarrow 1$ 
2  $s \leftarrow \sqrt{n}$ 
3 while  $i \leq s$  do
4    $a \leftarrow (i - 1)s + 1$ 
5    $b \leftarrow i \cdot s$ 
6    $P_i \leftarrow S_{a,b}$ 
7    $i \leftarrow i + 1$ 
8 end
```

Algorithm Transpose: Simulates a matrix transpose on the network

:

Input: A network of a collection of processor $P_1 \dots P_k$, m is the number of values that each processor keeps in memory ($m = \text{sqrtn}$ (see *Parallel-shear-sort*))

Result: If we think of M as a matrix where each processor P_i in M stores the i th row in M then after this routine returns each P_i will store the i th column of M

```
1 foreach  $P_i \in M$  do
2   for  $j \leftarrow 1$  to  $m$  do
3     if  $j \neq i$  then
4        $P_{j,i} \leftarrow P_{i,j}$ 
5     end
6   end
7 end
```

Algorithm Output-sequence: Output the sequence obtained by scanning over the network using snake like indices

:

Input: A network of a collection of proccessor $P_1 \dots P_k$, m is the number of values that each proccessor keeps in memory ($m = \text{sqrtn}$ see *Parallel-shear-sort*)

Output: Produces a sequence T from M according to the following rules. If i is even then scan P_i from left to right otherwise scan from right to left

```

1 foreach  $P_i \in M$  do
2    $T \leftarrow \text{empty-list}()$  if  $\text{is-even}(i)$  then
3     for  $j \leftarrow 1$  to  $m$  do
4        $T \leftarrow P_{i,j}$ 
5     end
6   else
7     for  $j \leftarrow m$  down-to 1 do
8        $T \leftarrow P_{i,j}$ 
9     end
10 end

```

Algorithm Parallel-sort: Produces a sorted version of S

:

Input: A network of a collection of proccessor $P_1 \dots P_k$, Sequence S , n is the number of values in S

```

1 Distribute-sequence( $M, S, n$ )
2  $\text{cycle\_num} \leftarrow 0$ 
3  $m \leftarrow \sqrt{n}$ 
4  $\text{max\_cycles} \leftarrow \log(m)$ 
5 while  $\text{cycle\_num} \leq \text{max\_cycles}$  do
6   if  $\text{is-even}(\text{cycle\_num})$  then
7     foreach  $P_i \in M$  do
8       if  $\text{is-even}(i)$  then
9          $\text{sort}(P_i, \text{increasing})$ 
10      else
11         $\text{sort}(P_i, \text{decreasing})$ 
12      end
13   else
14      $\text{Transpose}(M, m)$ 
15      $\text{sort}(P_i, \text{increasing})$ 
16      $\text{Transpose}(M, m)$ 
17      $\text{cycle\_num} \leftarrow \text{cycle\_num} + 1$ 
18 end
19 Output-sequence( $M, m$ )

```

1.2 Analysis

1.2.1 Amount of compare exchange operations performed by each processor

Recall that we use merge sort to perform compare exchange operations. Thus, the total number of comparisons performed between all the processors in a given cycle is $O(n \log n)$. We know by the theorem stated in the slides that the number of iterations until *Parallel-shear-sort* converges is at most $O(\log \sqrt{n}) = O(\frac{1}{2} \log(n))$. This means the total number of compare exchange operations performed over *all* cycles is at most $(\frac{1}{2} \log n)(n \log n) = \frac{n}{2} \log^2(n)$

1.2.2 Amount of memory needed by each processor

Each processor needs to be able to store \sqrt{n} items. In total $O(n)$ storage is needed.

1.2.3 Number of Cross-Bar communications cycles

The only time communications between processors takes place is during the transpose phase of *Parallel-shear-sort* in lines 13, 15. This means only $\frac{1}{2}$ the time (during an odd cycle) communications will take place.

Transpose is invoked twice in the else block. Each invocation requires n communications which mean there will be $2n$ during an odd cycle.

In 1.2.1 we showed that there are no more than $\frac{1}{2} \log n$ cycles.

Putting all these ideas together, we see the total number of cross bar communications performed by the algorithm is ...

$$(\frac{1}{2})(2n)(\frac{1}{2} \log n) = \frac{n \log n}{2}$$

which is $O(n \log n)$

2 Bitonic Sort

I assume the $n = 2^j$ for some positive integer j and that n is divisible by k , the number of processor I assume the existence of two function, *send* and *read* responsible for communication between processors on the network. These functions are not defined explicitly but the headers are given. Also, note that *Compare-exchange* is defined to work when communication is between two different processors or just within the same processor. *Output-sequence* here simply scans over each processors memory starting at lower index processors and going up outputting the contents of there memory. Since it is such a simple function I have not included it. Lastly, I assume that the input sequence S is a bitonic sequence where the first half of the sequence is nondecreasing and the second half nonincreasing. This assumption was permitted by the instructor. Futhermore, even if it wasnt, it wouldnt be very hard to transform an arbitrary sequence so that it was bitonic

2.1 Implementation

Algorithm send: Sends data from one processor to another

Input: The sending procceesor P_i , the element to be sent x , the receiving proccesor P_j

Result: Element x is sent from P_i to P_j . Whatever is in P_j read buffer is automatically overwritten with x

Algorithm read: procceesor P_i reads its read buffer and returns whatever value is in there.

Input: The processor whose read buffer is to be read

Result: The value inside of P_i read buffer. There is only a valid value here if something was sent to P_i from another processor. Otherwise the buffer is garbage

Algorithm Distribute-sequence: Stores a subsequence of S in the processors

Input: A network of a collection of processor $P_1 \dots P_k$, Sequence S , n is number of values in S , k is the number of processors

Result: Calculates how many values to load into each processor based on n then reads that many consecutive values into the correct processor

```
1  $i \leftarrow 1$ 
2  $s \leftarrow \frac{n}{k}$ 
3 while  $i \leq k$  do
4    $a \leftarrow (i - 1)(s) + 1$ 
5    $b \leftarrow i \cdot s$ 
6    $P_i \leftarrow S_{a,b}$ 
7    $i \leftarrow i + 1$ 
8 end
```

Algorithm Compare-exchange: Compares the elements between two processors and swaps them if the first processors element is greater than the seconds

Input: A processor P_i , the index of the element for comparison in P_i called a , A second processor P_j , the index of the element for comparison in P_j called b

Result: We assume that P_i should store smaller values and P_j larger. Thus, P_i and P_j send eachother values (if they are different proccesors) and P_i keeps the minimum of the two values whereas P_j keeps the maximum of the two

```
1 if  $i \neq j$  then
2    $\text{send}(P_i, P_{i,a}, P_j)$ 
3    $\text{send}(P_j, P_{j,b}, P_i)$ 
4    $x \leftarrow \text{read}(P_i)$ 
5    $P_{i,a} \leftarrow \min(P_{i,a}, x)$ 
6    $x \leftarrow \text{read}(P_j)$ 
7    $P_{j,b} \leftarrow \min(P_{j,b}, x)$ 
8 else
9    $x \leftarrow P_{i,a}$ 
10  if  $x > P_{i,b}$  then
11     $P_{i,a} \leftarrow P_{i,b}$ 
12     $P_{i,b} \leftarrow x$ 
13  end
```

Algorithm Single-processor-bitonic-sort: Given a proccessor P whose memory stores a bitonic sequence.

:

Input: A single proccessor from the network P_i , n number of proccessors in network that P_i is from, k number proccessor in network that P_i is from

```

1  $list\_size \leftarrow \frac{n}{k}$ 
2  $j \leftarrow 1$ 
3  $offset \leftarrow list\_size \cdot 2^{-j}$ 
4 while  $offset \geq 1$  do
5   for  $i = 1$  to  $(list\_size \cdot 2^{-1})$  do
6      $\mid$  Compare-exchange( $P_i, i, P_i, i + offset$ )
7   end
8 end

```

Algorithm Parallel-bitonic-sort: Produces a sorted sequence given a bitonic sequence S where the first half of S is nondecreasing and the second half is nonincreasing

:

Input: A network of a collection of proccessor $P_1 \dots P_k$, Bitonic sequence S , n is number of values in S , k is the number of proccessors

```

1 Distribute-sequence( $M, S, n, k$ )
2  $list\_size \leftarrow \frac{n}{k}$ 
3  $j \leftarrow 1$ 
4  $offset \leftarrow list\_size \cdot 2^{-j}$ 
5 while  $offset \geq 1$  do
6   for  $i = 1$  to  $\frac{k}{2}$  do
7     for  $c = 1$  to  $(list\_size \cdot 2^{-1})$  do
8        $\mid$   $a \leftarrow i$ 
9        $\mid$   $b \leftarrow i + offset$ 
10       $\mid$  Compare-exchange( $P_a, c, P_b, c$ )
11     end
12   end
13 end
14 foreach  $P_i$  do
15    $\mid$  Single-procceeisor-bitoncic-sort( $P_i$ )
16 end
17 Output-sequence( $M$ )

```

2.2 Analysis

2.2.1 Amount of compare exchange operations performed by each processor

There will be $\frac{n}{4} \log n$ iterations for the while loop in *parallel-bitonic-sort*. *Single-processor-bitonic-sort* is called once for each processor after the while loop. There are $n \log n$ total comparisons as a result of this function. Thus in total there are $O(n \log n)$ communications amongst processors.

2.2.2 Amount of memory needed by each processor

A given processor will need to store $\frac{n}{k}$ elements. In total of course the storage is $O(n)$

2.2.3 Number of Cross-Bar communications cycles

In this case the analysis of the cross bar communication is identical to 2.2.1 thus there are $O(n \log n)$ cross bar communications cycles